

Serializability

- ▶ A schedule is serializable if it is **equivalent to a serial schedule**.
- ▶ In **serial schedules**, only **one transaction is allowed to execute at a time** i.e. **no concurrency is allowed**.
- ▶ Whereas in **serializable schedules**, **multiple transactions can execute simultaneously** i.e. **concurrency is allowed**.
- ▶ Types (forms) of serializability
 - Conflict serializability
 - View serializability

Conflicting instructions

► Let I_i and I_j be two instructions of transactions T_i and T_j respectively.

1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$
 I_i and I_j don't conflict

T_i	T_j
read (Q)	
	read (Q)

T_i	T_j
	read (Q)
read (Q)	

2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$
 I_i and I_j conflict

T_i	T_j
read (Q)	
	write(Q)

T_i	T_j
	read (Q)
write(Q)	

3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$
 I_i and I_j conflict

T_i	T_j
write(Q)	
	read (Q)

T_i	T_j
	write(Q)
read (Q)	

4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$
 I_i and I_j conflict

T_i	T_j
write(Q)	
	write(Q)

T_i	T_j
	write(Q)
write(Q)	

Conflict serializability

- If a given schedule can be **converted into a serial schedule by swapping its non-conflicting operations**, then it is called as a conflict serializable schedule.

T1	T2
Read (A) A = A - 50 Write (A)	
	Read (A) Temp = A * 0.1 A = A - temp Write (A)
Read (B) B = B + 50 Write (B) Commit	
	Read (B) B = B + temp Write (B) Commit

T1	T2
Read (A) A = A - 50 Write (A) Read (B) B = B + 50 Write (B) Commit	
	Read (A) Temp = A * 0.1 A = A - temp Write (A) Read (B) B = B + temp Write (B) Commit

Conflict serializability (Example)

- ▶ Example of a **schedule that is not conflict serializable**:

T1	T2
Read (A)	Write (A)
Read (A)	

- ▶ We are **unable to swap instructions** in the above schedule to obtain either the serial schedule $\langle T1, T2 \rangle$, or the serial schedule $\langle T2, T1 \rangle$.




Cycle is formed, i.e., it is not conflict serializable schedule

Conflict serializability (Example)

► Is given **schedule conflict serializable**?

T1	T2	T3
Read (A)		Read (Z) Write (Z)
Read (Y)	Read (Y) Write (Y)	
Write (X)	Write (Z)	Write (X)




Cycle is formed, i.e., it is not conflict serializable schedule

Conflict serializability (Example)

► Is given **schedule conflict serializable**?

T1	T2	T3
Read (X)	Read (Y)	Read (Y)
Write (X)	Write (Y)	Write (X)
	Read (X)	
	Write (X)	




Cycle is not formed, i.e., it is a conflict serializable schedule

Conflict serializability (Example)

► Is given **schedule conflict serializable**?

T1	T2	T3
Read (a)	Read (b)	Read (c) Write (c)
Write (a)	Write (b)	




Cycle is not formed, i.e., it is a conflict serializable schedule

Conflict serializability (Example)

► Is given **schedule conflict serializable**?

T1	T2	T3	T4
Write (X) Commit	Read (X) Write (Y) Read (Z) Commit	Write (X) Commit	Read (X) Read (Y) Commit




Cycle is not formed, i.e., it is a conflict serializable schedule

Conflict serializability (Example)

► Is given **schedule conflict serializable**?

T1	T2	T3
Read (a)	Write (a)	
Write (a)		
		Write (a)



Cycle is formed, i.e., it is not conflict serializable schedule

View serializability

- ▶ Let S_1 and S_2 be two schedules with the same set of transactions. S_1 and S_2 are view equivalent if the following three conditions are satisfied, for each data item Q
 - ➔ Initial Read
 - ➔ Updated Read
 - ➔ Final Write
- ▶ If a schedule is view equivalent to its serial schedule then the given schedule is said to be view serializable.

Initial Read

- ▶ If in **schedule S1**, transaction **T_i** reads the **initial value of Q**, then in **schedule S2** also **transaction T_i** must read the **initial value of Q**.

S1	
T1	T2
Read (A)	Write (A)

S2	
T1	T2
Read (A)	Write (A)

S3	
T1	T2
Write (A)	Read (A)

- ▶ Above two schedules **S1** and **S3** are **not view equivalent** because **initial read operation in S1 is done by T1** and in **S3** it is done by **T2**.
- ▶ Above two schedules **S1** and **S2** are **view equivalent** because **initial read operation in S1 is done by T1** and in **S2** it is also done by **T1**.

Updated Read

- ▶ If in **schedule S1**, transaction T_i executes $\text{read}(Q)$, and that value was produced by transaction T_j (if any), then in **schedule S2** also transaction T_i must read the value of Q that was produced by transaction T_j .

S1		
T1	T2	T3
Write (A)	Write (A)	Read (A)

S2		
T1	T2	T3
Write (A)	Write (A)	Read (A)

S3		
T1	T2	T3
Write (A)	Write (A)	Read (A)

- ▶ Above two schedules **S1** and **S3** are not view equal because, in **S1**, T_3 is reading A that is updated by T_2 and in **S3**, T_3 is reading A which is updated by T_1 .
- ▶ Above two schedules **S1** and **S2** are view equal because, in **S1**, T_3 is reading A that is updated by T_2 and in **S2** also, T_3 is reading A which is updated by T_2 .

Final Write

- ▶ If T_i performs the final write on the data value in $S1$, then it also performs the final write on the data value in $S2$.

S1		
T1	T2	T3
Write (A)	Read (A)	Write (A)

S2		
T1	T2	T3
Write (A)	Read (A)	Write (A)

S3		
T1	T2	T3
Write (A)	Write (A)	Read (A)

- ▶ Above two schedules **S1 and S3 are not view equal** because **final write operation in S1 is done by T3 and in S3 final write operation is also done by T1**.
- ▶ Above two schedules **S1 and S2 are view equal** because **final write operation in S1 is done by T3 and in S2 also the final write operation is also done by T3**.