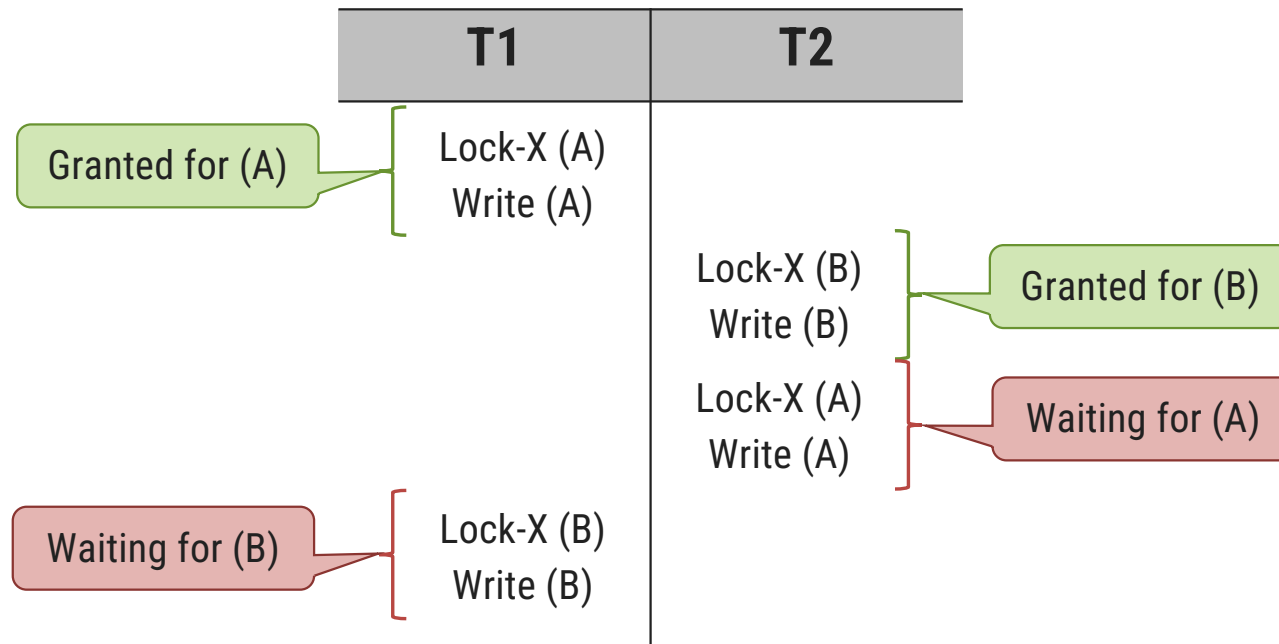


Deadlock

What is deadlock?

- Consider the following two transactions:




- A deadlock is a **situation in which two or more transactions are waiting for one another to give up locks.**

Deadlock detection

- ▶ A simple way to detect deadlock is with the help of **wait-for graph**.
- ▶ One **node is created** in the wait-for graph for **each transaction that is currently executing**.
- ▶ Whenever a **transaction T_i is waiting to lock an item X that is currently locked by a transaction T_j , a directed edge from T_i to T_j ($T_i \rightarrow T_j$) is created in the wait-for graph**.
- ▶ When **T_j releases the lock(s) on the items that T_i was waiting for**, the **directed edge is dropped** from the wait-for graph.
- ▶ We have a state of **deadlock if and only if the wait-for graph has a cycle**.
- ▶ Then **each transaction involved in the cycle is said to be deadlocked**.


Deadlock detection

- ▶ Transaction **A is waiting for** transactions **B and C**.
- ▶ Transaction **C is waiting** for transaction **B**.
- ▶ Transaction **B is waiting** for transaction **D**.
- ▶ This wait-for graph has **no cycle**, so there is **no deadlock state**.
- ▶ Suppose now that transaction **D is requesting an item held by C**. Then the **edge $D \rightarrow C$** is added to the wait-for graph.
- ▶ Now this **graph contains the cycle**.
- ▶ $B \rightarrow D \rightarrow C \rightarrow B$
- ▶ It means that **transactions B, D and C are all deadlocked**.



Deadlock recovery

- ▶ When a deadlock is detected, the system must recover from the deadlock.
- ▶ The most common **solution is to roll back one or more transactions to break the deadlock.**
- ▶ Choosing which transaction to abort is known as **victim selection.**
- ▶ In this wait-for graph transactions B, D and C are deadlocked.
- ▶ In order to remove deadlock one of the transaction out of these three (B, D, C) transactions must be roll backed.
- ▶ We should **rollback those transactions that will incur the minimum cost.**
- ▶ When a deadlock is detected, the choice of which transaction to abort can be made using following criteria:
 - The transaction which have the fewest locks
 - The transaction that has done the least work
 - The transaction that is farthest from completion



Deadlock prevention

- ▶ A protocols **ensure that the system will never enter into a deadlock state.**
- ▶ Some prevention strategies :
 - ↪ Require that **each transaction locks all its data items before it begins execution** (pre-declaration).
 - ↪ Impose **partial ordering on all data items and a transaction can lock data items only in the specified order.**
- ▶ Drawbacks:
 - ↪ It is difficult to predict what data items are required.
 - ↪ Data item utilization will be low.
 - ↪ Ordering of the data item may be difficult and time consuming.

Deadlock prevention

► Following schemes use transaction timestamps for the sake of deadlock prevention alone.

1. Wait-die scheme – non-preemptive

- If an **older transaction is requesting a resource** which is held by younger transaction, then **older transaction is allowed to wait** for it till it is available.
- If an **younger transaction is requesting a resource** which is held by older transaction, then **younger transaction is killed**.

2. Wound-wait scheme – preemptive

- If an **older transaction is requesting a resource** which is held by younger transaction, then **older transaction forces younger transaction to kill** the transaction and release the resource.
- If an **younger transaction is requesting a resource** which is held by older transaction, then **younger transaction is allowed to wait** till older transaction will releases it.

	Wait-die	Wound-wait
O needs a resource held by Y	O waits	Y dies
Y needs a resource held by O	Y dies	Y waits

Deadlock prevention

▶ Following schemes use transaction timestamps for the sake of deadlock prevention alone.

3. Timeout-Based Schemes

- ↪ A **transaction waits for a lock only for a specified amount of time. After that, the wait times out and the transaction is rolled back.** So deadlocks never occur.
- ↪ **Simple to implement; but difficult to determine good value of the timeout interval.**